

The background features a vibrant blue gradient with subtle, wavy horizontal lines. In the bottom right corner, there are abstract, flowing shapes in shades of purple, pink, and orange, creating a dynamic and modern aesthetic.

aws SUMMIT

INDIA | MAY 25, 2023

B M A R C 0 0 1

Enterprise Architectures - State of the Union

Mani Chandrasekaran

Principal Solutions Architect – India and South Asia
AWS India



Agenda

- The value of architecture patterns
- Common integration patterns
- Application patterns – Serverless and Containers
- What next?

AWS Global Infrastructure



Why customers choose AWS

Most experience

16

years helping millions
of customers

Global reach & high availability

99

availability zones spanning 31
geographic regions

Security & compliance

300+

security
features

**Customer obsession
& innovation**

200+

service offerings

**AWS's Infrastructure
is**

3x

More Energy Efficient

than the median of surveyed U.S.
enterprise data centers

**Improve
TCO**

111

price reductions since 2006

Machine learning

81%

of all deep learning is running on AWS¹

Ecosystem

4,500

software listings from 1,400 ISVs



AWS Global Infrastructure

31 Launched Regions, 99 Availability Zones and 450+ Points of Presence



United States

GovCloud (U.S.):

U.S.-East (3), US-West (3)

U.S. West

Oregon (4), Northern California (3)

U.S. East

N. Virginia (6), Ohio (3)



Canada

Central (3)



South America

São Paulo (3)



Europe

Frankfurt (3)

Ireland (3)

London (3)

Milan (3)

Paris (3)

Spain (3)

Stockholm (3)

Zurich (3)



Africa

Cape Town (3)



Middle East

Bahrain (3)

UAE (3)



Asia Pacific

*Beijing, operated by Sinnet (3)

*Ningxia, operated by NWCD (3)

Hong Kong (3)

Hyderabad (3)

Jakarta (3)

Mumbai (3)

Osaka (3)

Seoul (4)

Singapore (3)

Tokyo (4)



Australia

Sydney (3)

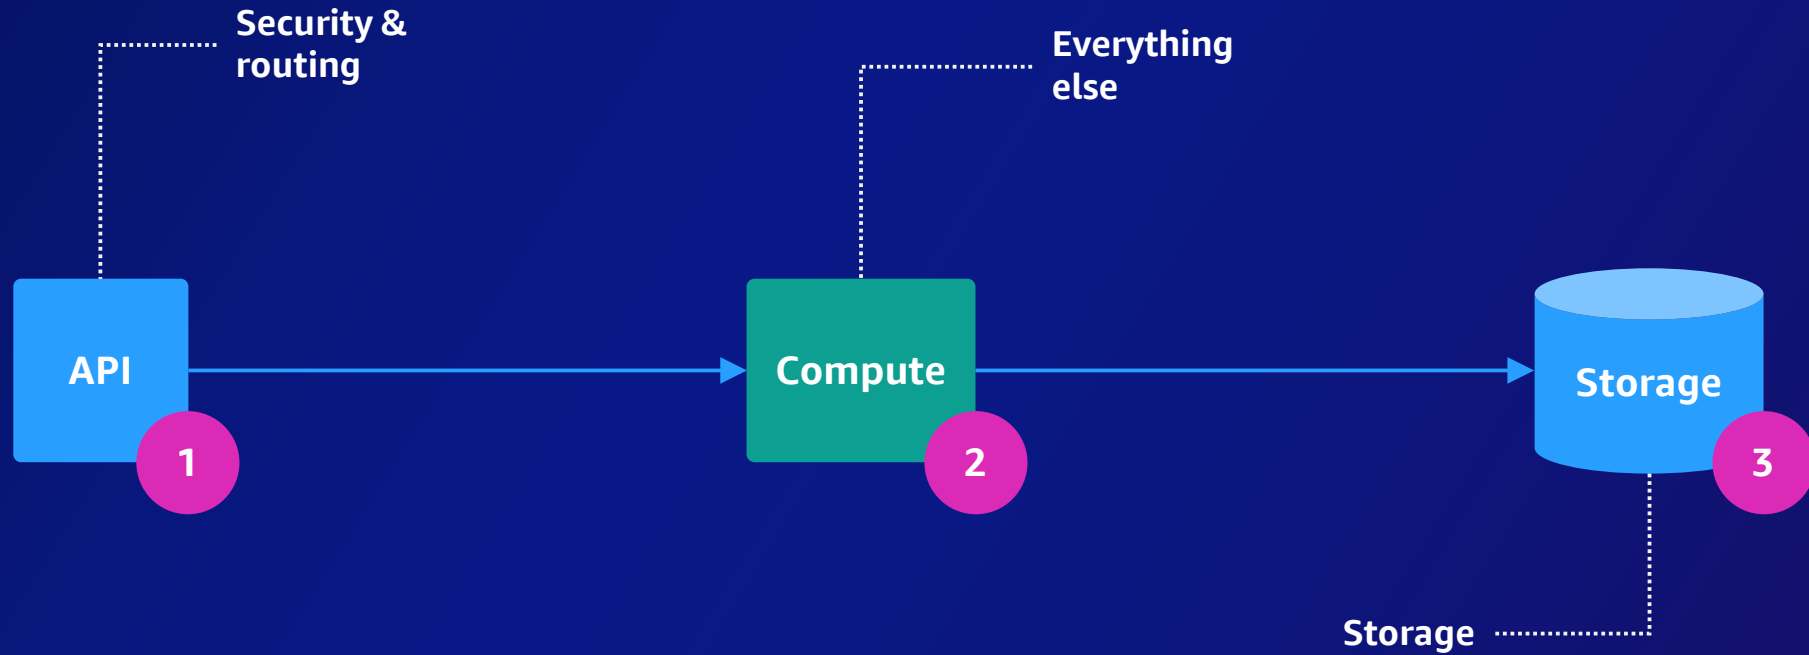
Melbourne (3)

Latest Region & number of Availability Zones (AZs) can be found here [AWS global infrastructure page](#)

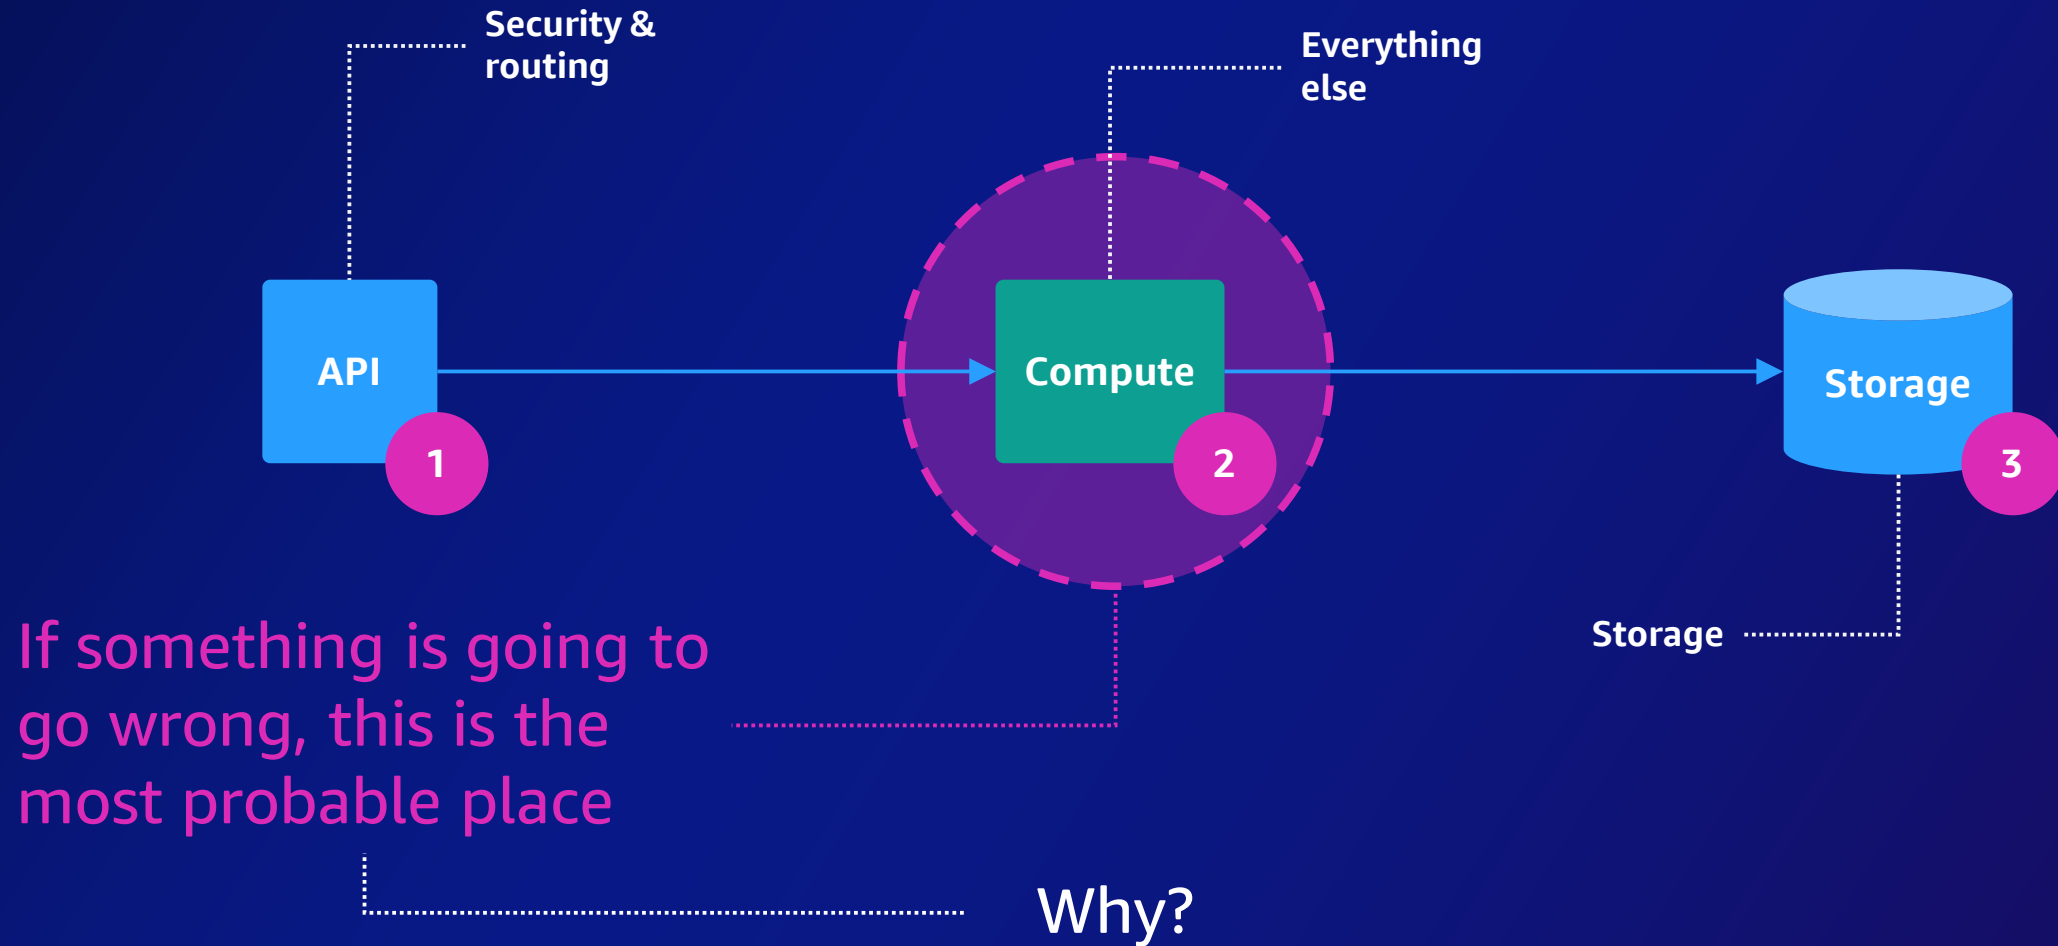


Architecture patterns in enterprise architectures

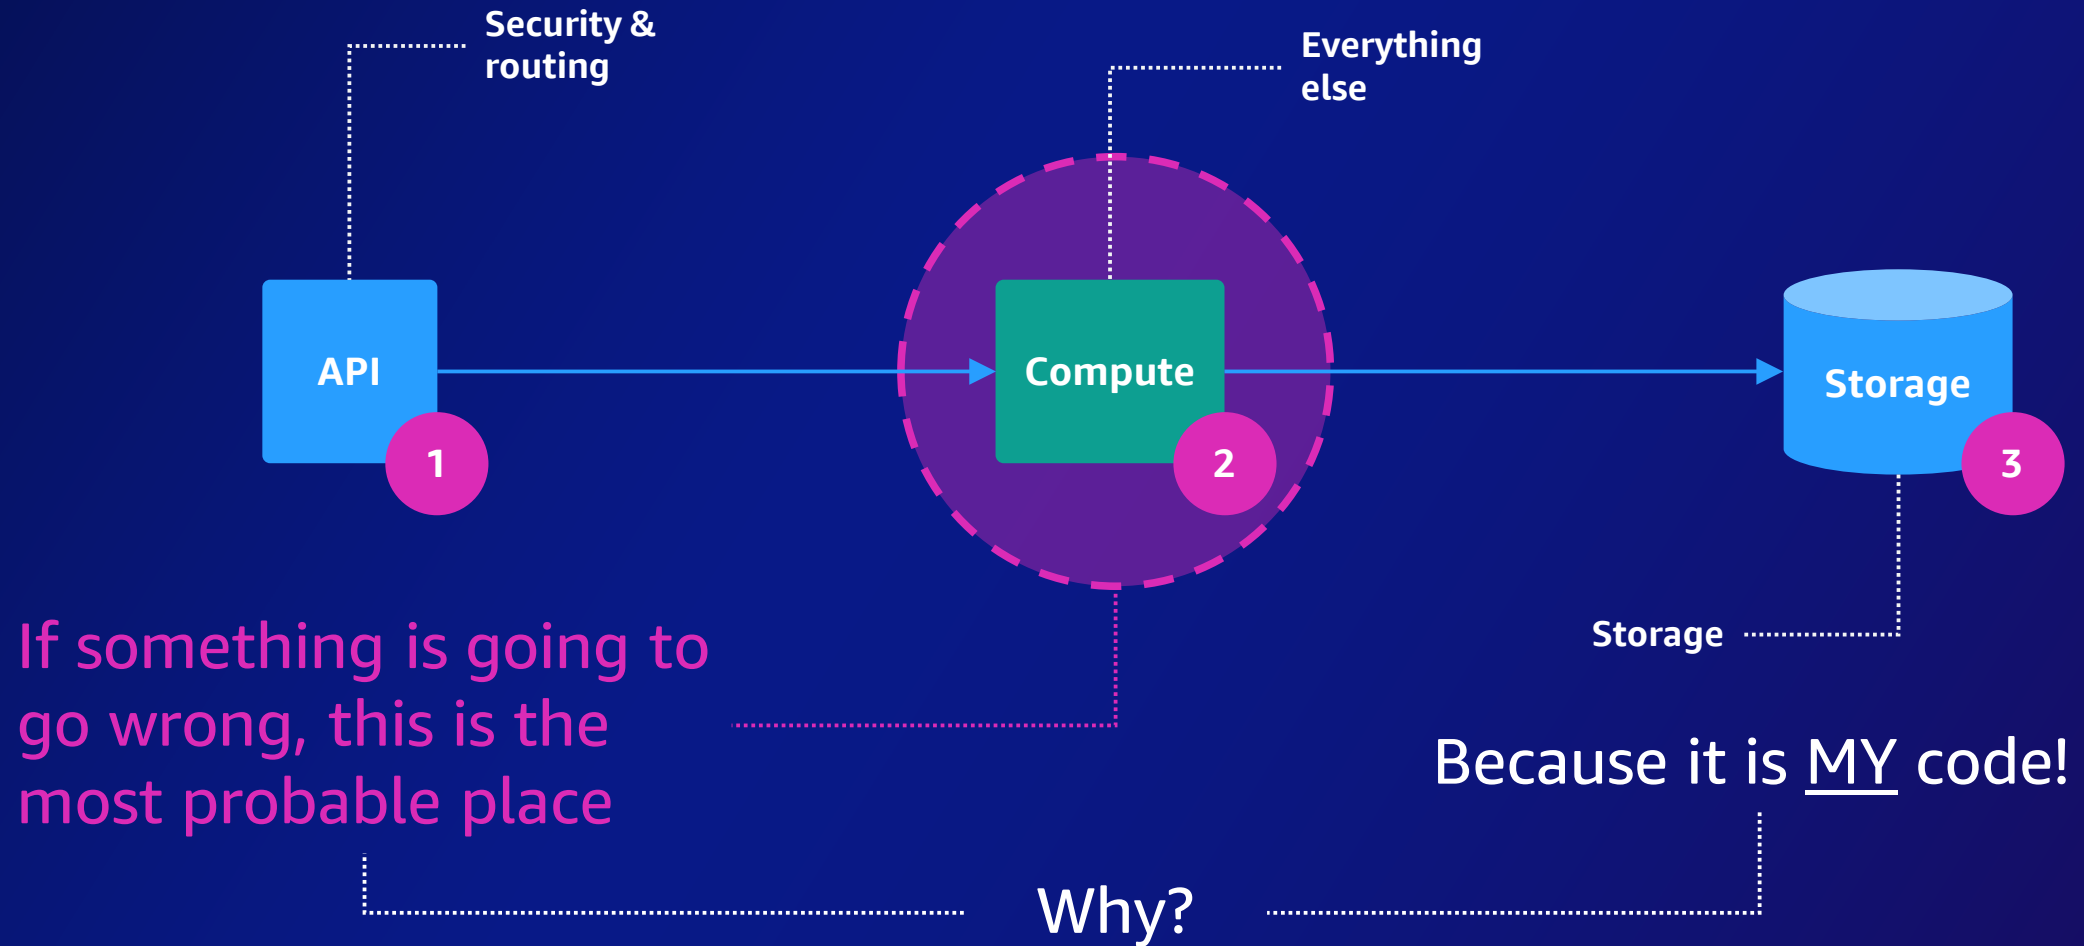
Application elements




Application elements



Application elements



How do we improve this?



Synchronous systems

- Are inherently tightly coupled
- Problems downstream can immediately impact upstream
- Retries from upstream callers can amplify problems



Synchronous systems

- Are inherently **tightly** coupled
- Problems don't propagate immediately
- Retries from clients amplify problems

...and some things simply
take **too much** time to **wait**,
or are **asynchronous** by **nature**!

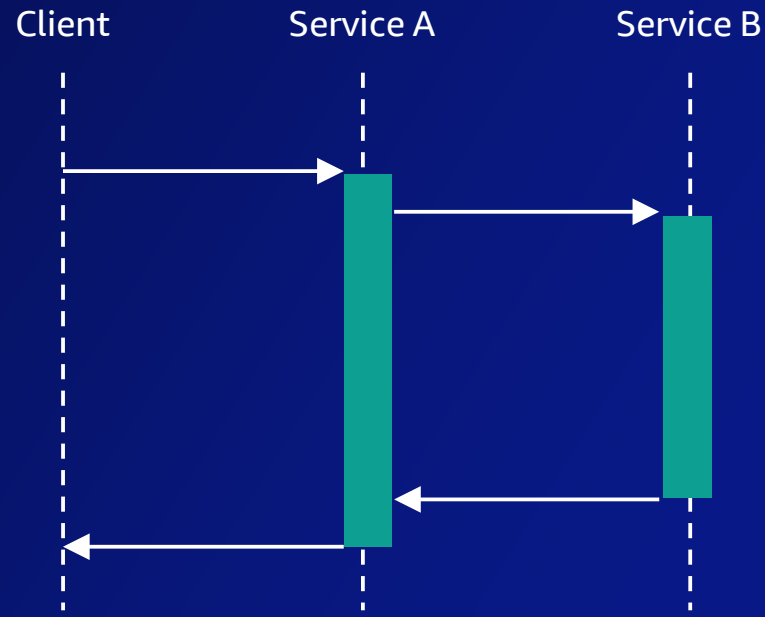
“ If your application is cloud-native, or large-scale, or distributed, and doesn’t include a messaging component, that’s probably a bug ”

Tim Bray


General-purpose internet-software geek

Thinking asynchronously

Think more asynchronously



**Synchronous
commands**




**Asynchronous
events**

If you don't need a response, execute asynchronously

Synchronous



Asynchronous



Integration patterns


In modern cloud applications, **integration** isn't an afterthought. It's an **integral part** of the **application architecture** and the software delivery lifecycle.

Gregor Hohpe

Author of [Enterprise Integration Patterns](#), [Cloud Strategy](#), and [The Software Architect Elevator](#)

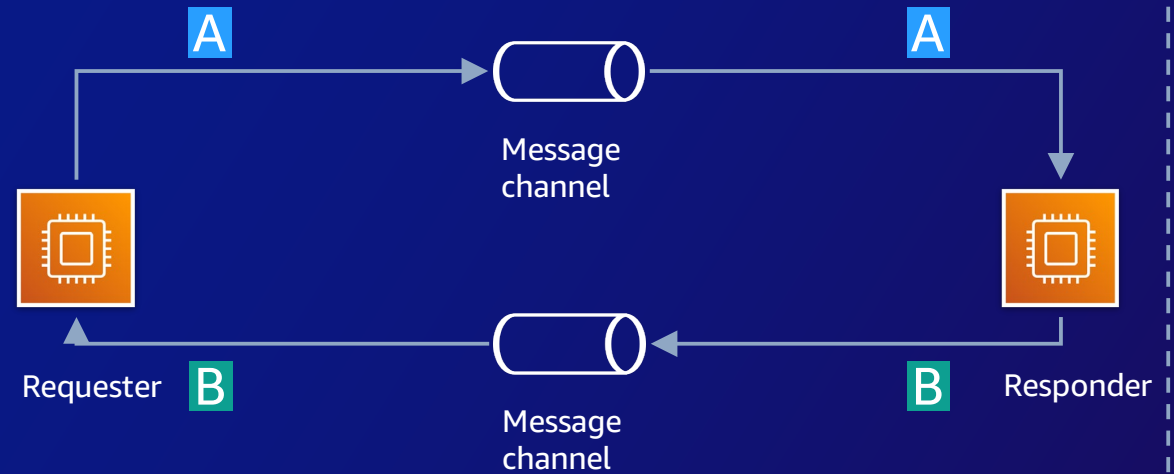
Message exchange

One-way



No response expected

Request-response



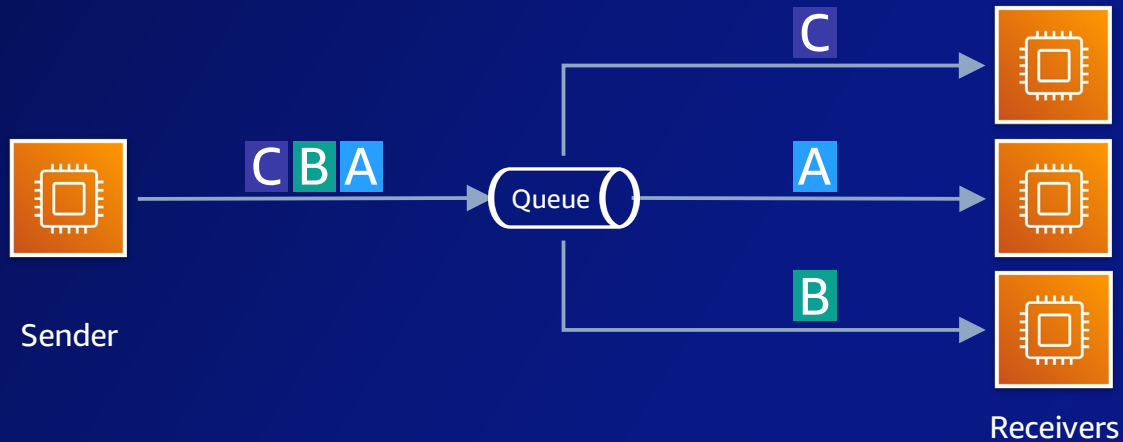
Response expected

Return address

Correlation ID

Message channels

Point-to-point (queue)

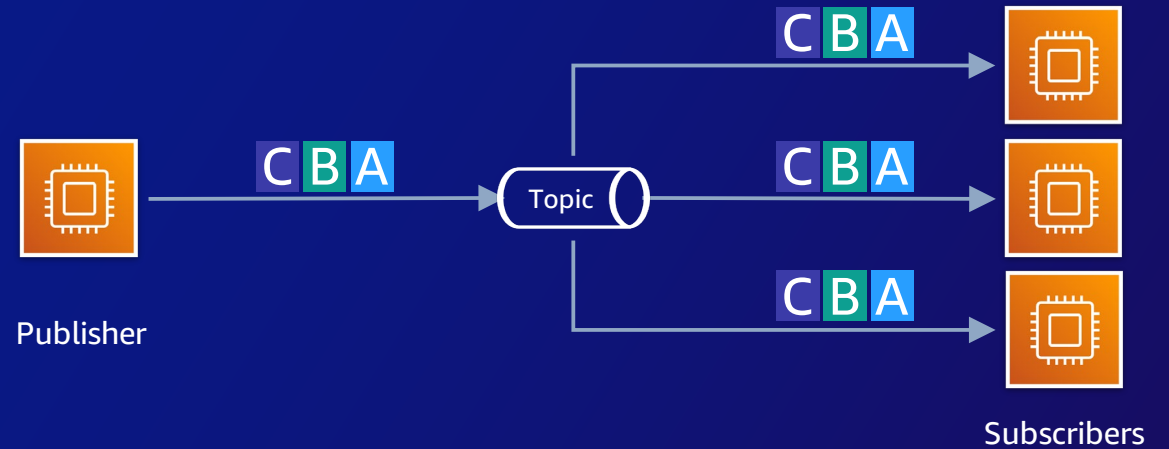


Consumed by one receiver

Easy to scale

Flatten peak loads

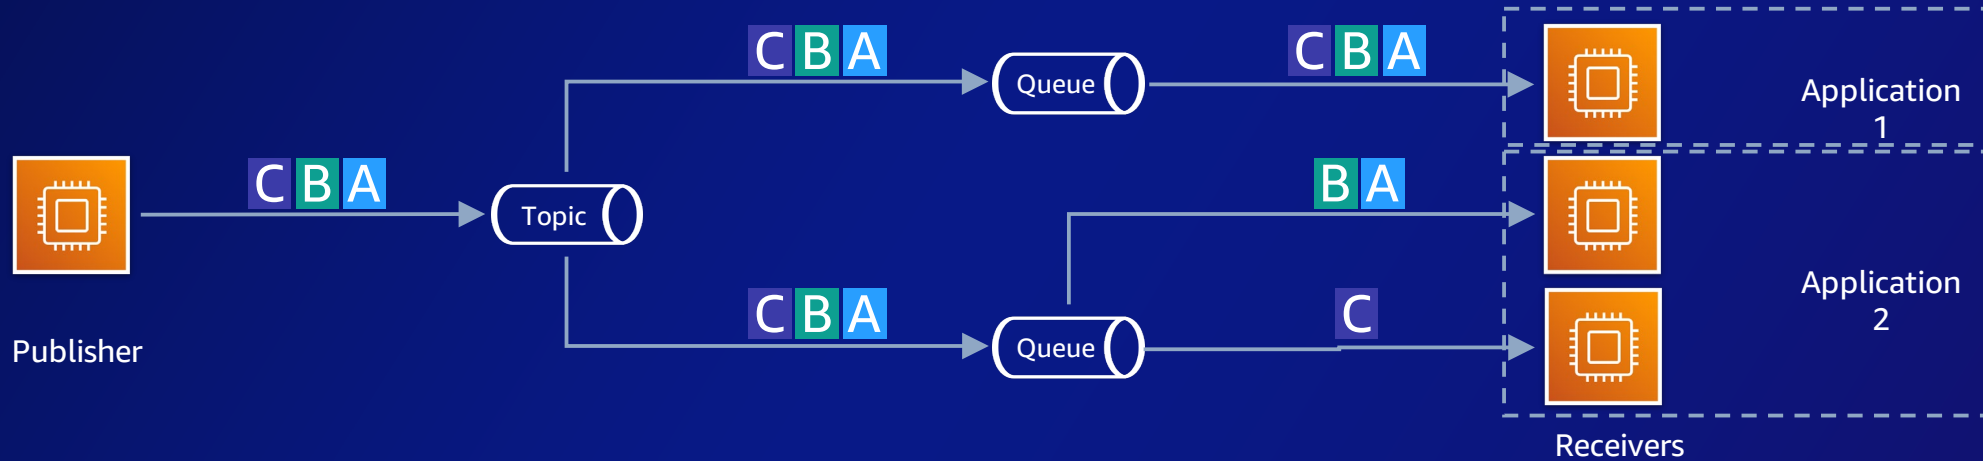
Publish-subscribe (topic)



Consumed by all subscribers

Message channels

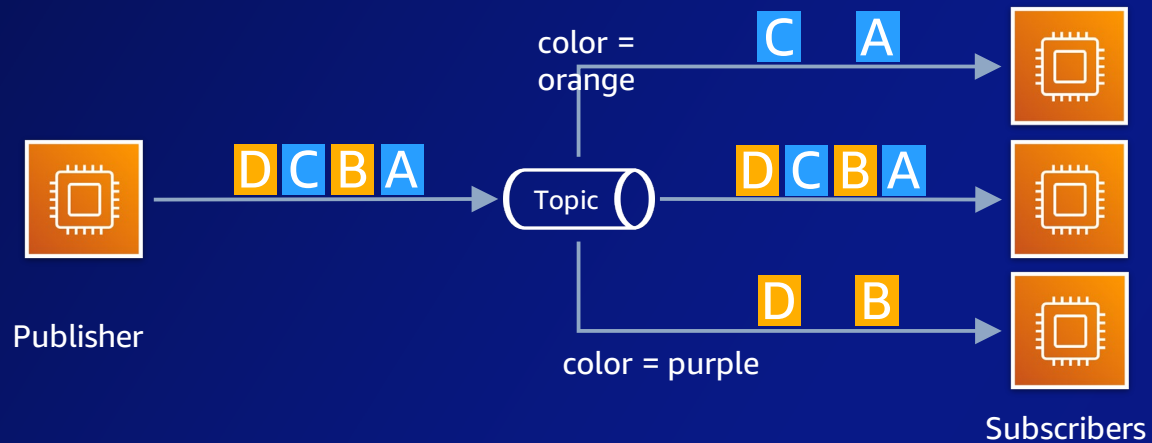
Topic-queue-chaining



Allows fan-out and receiver scale-out at the same time

Message routing

Message filter

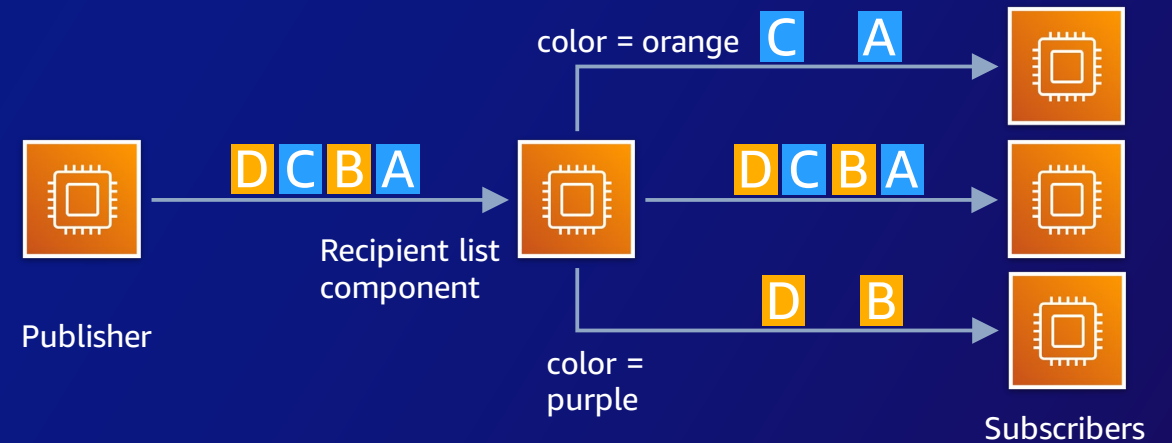


Receive only a relevant subset of messages

Controlled by subscriber

Publisher remains completely unaware

Recipient list



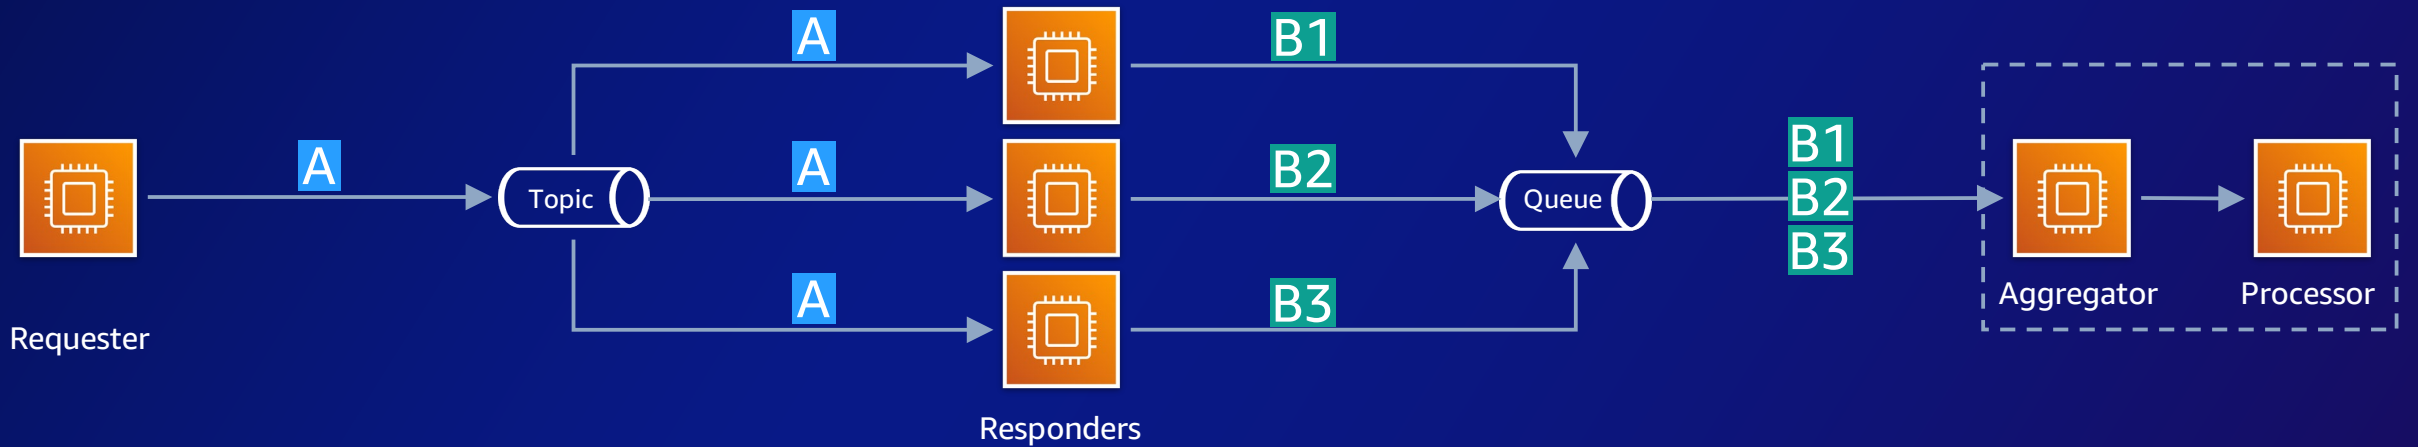
Send only a relevant subset of messages to a subscriber

Controlled by publisher or separate component

Potentially adds coupling

Message routing

Scatter-gather



How to distribute a request across potentially interested/relevant parties and capture their individual responses?

Election or parallel processing scenarios, i.e., search for **best** response or **accumulate** responses

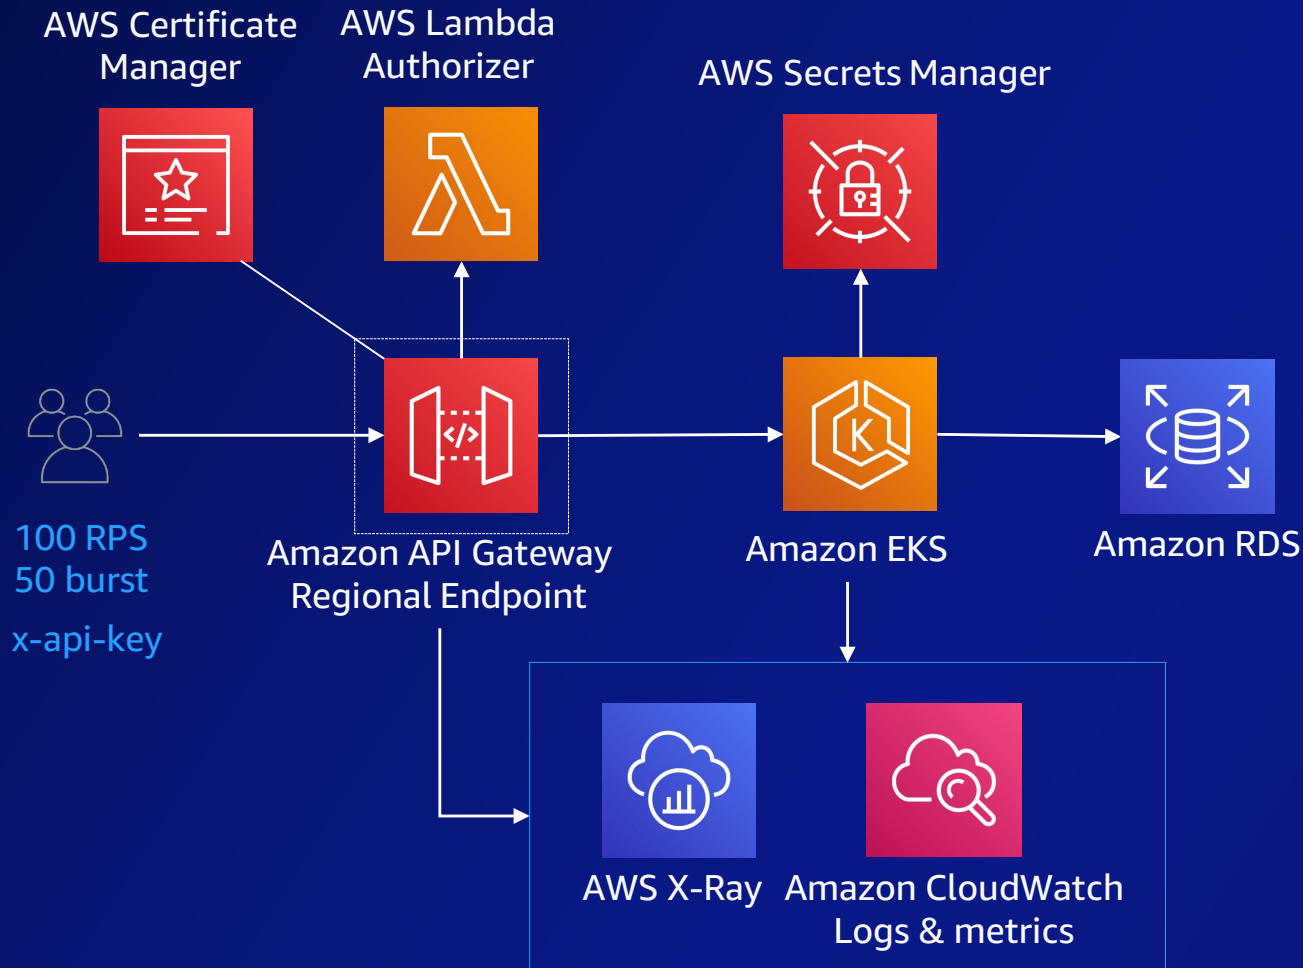
Application patterns – Serverless and Containers

REST API pattern

REST API



REST API with Containers



- ❑ Enable access logs, execution logs for your APIs
- ❑ API Gateway is well integrated with AWS X-Ray and CloudWatch
- ❑ Instrument your backend services to capture X-Ray traces and create metrics with CloudWatch Embedded Metric Format or Prometheus Exposition Format
- ❑ Regulate inbound access rates
- ❑ Authorize consumers. Enforce API keys
- ❑ Secure end points with SSL certificates
- ❑ Manage secrets with AWS Secrets Manager
- ❑ Regional endpoints support HTTP2

REST API with AWS Lambda



Best practices

- Enable access logs, structure logs and instrument your code
- Create metrics async with CloudWatch Embedded Metric Format
- Regulate inbound access rates
- Authorize consumers. Manage secrets with AWS Secrets Manager
- On-demand tables support up to 40K read/write request units
- Regional endpoints support HTTP2

OPERATIONS

RELIABILITY

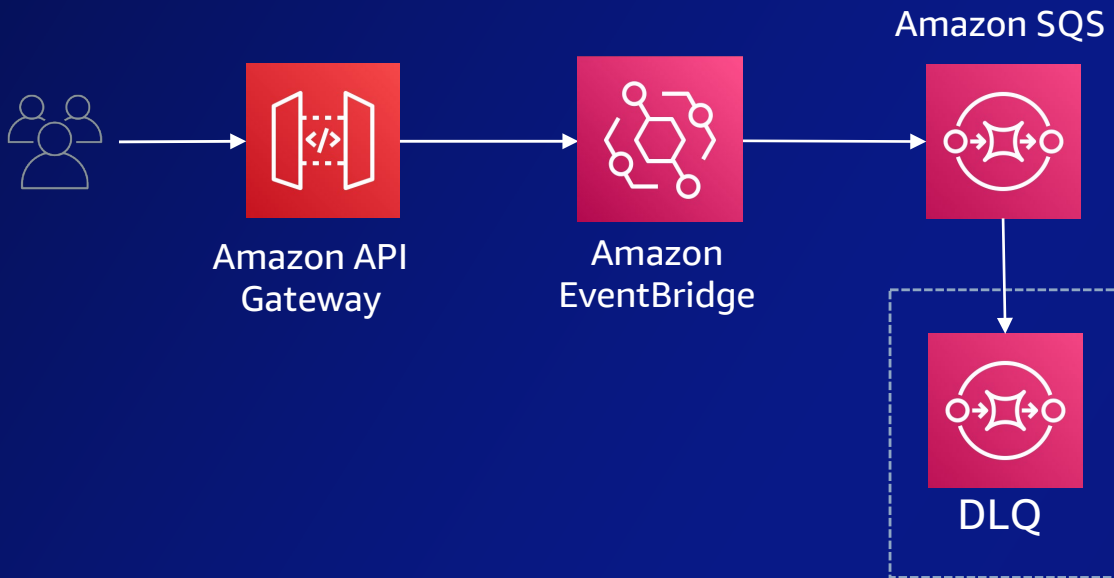
SECURITY

PERFORMANCE

COST

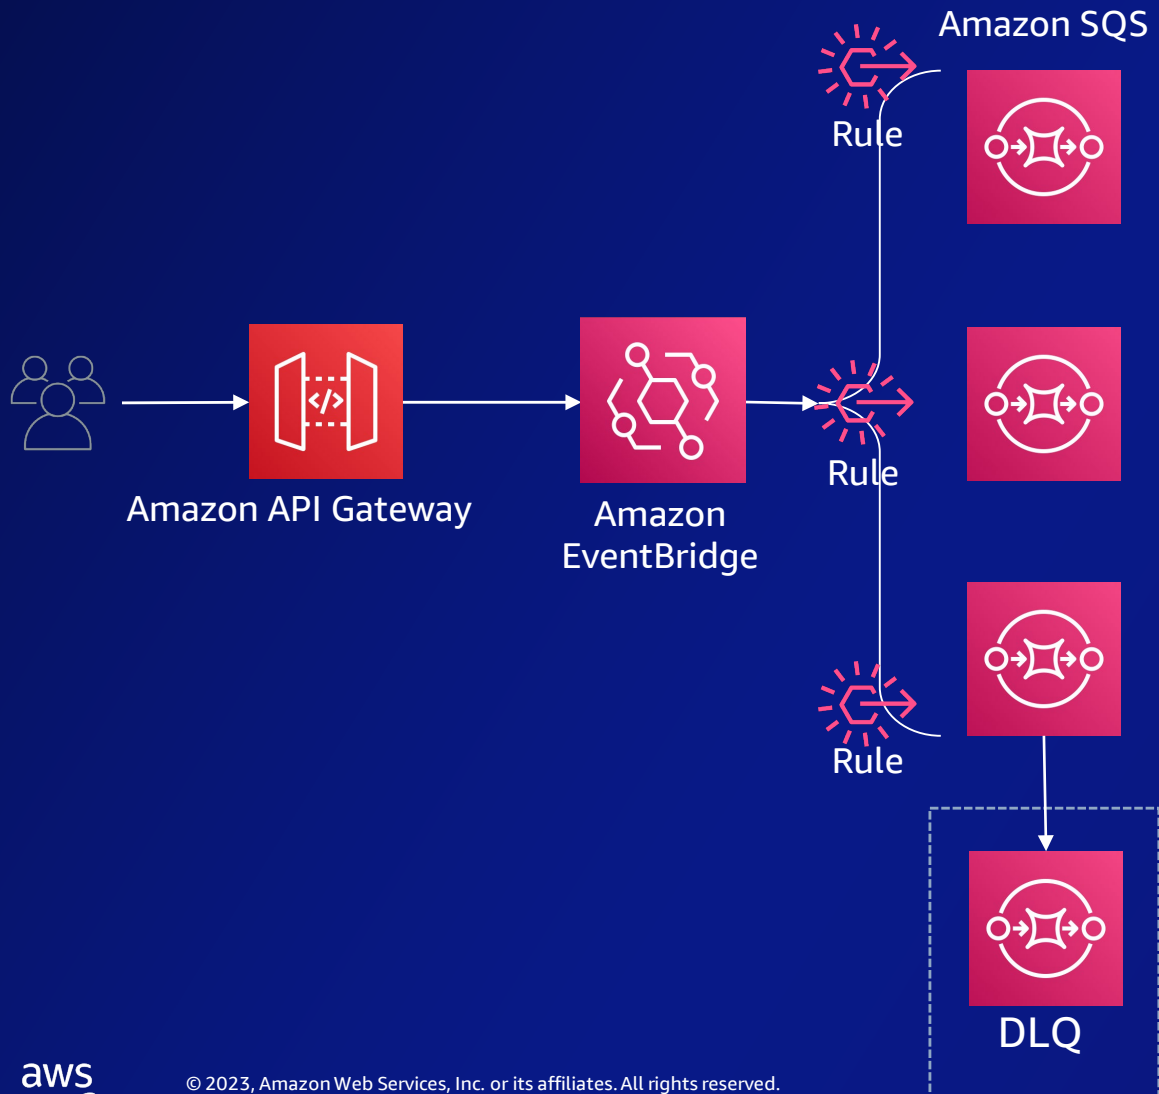
Fan-out pattern

Fan-out pattern for message processing



- ❑ API Gateway can integrate with AWS services directly
- ❑ EventBridge is a pub/sub service to decouple producer and consumer
- ❑ Publish notifications directly to EventBridge from API Gateway
- ❑ Integrate with Amazon SQS for higher durability, batching, and DLQ

Fan-out pattern for message processing



- ❑ API Gateway can integrate with AWS services directly
- ❑ Publish notifications directly to EventBridge
- ❑ Integrate with Amazon SQS for higher durability, batching, and DLQ
- ❑ Use EventBridge rules to separate messages into different SQS queues

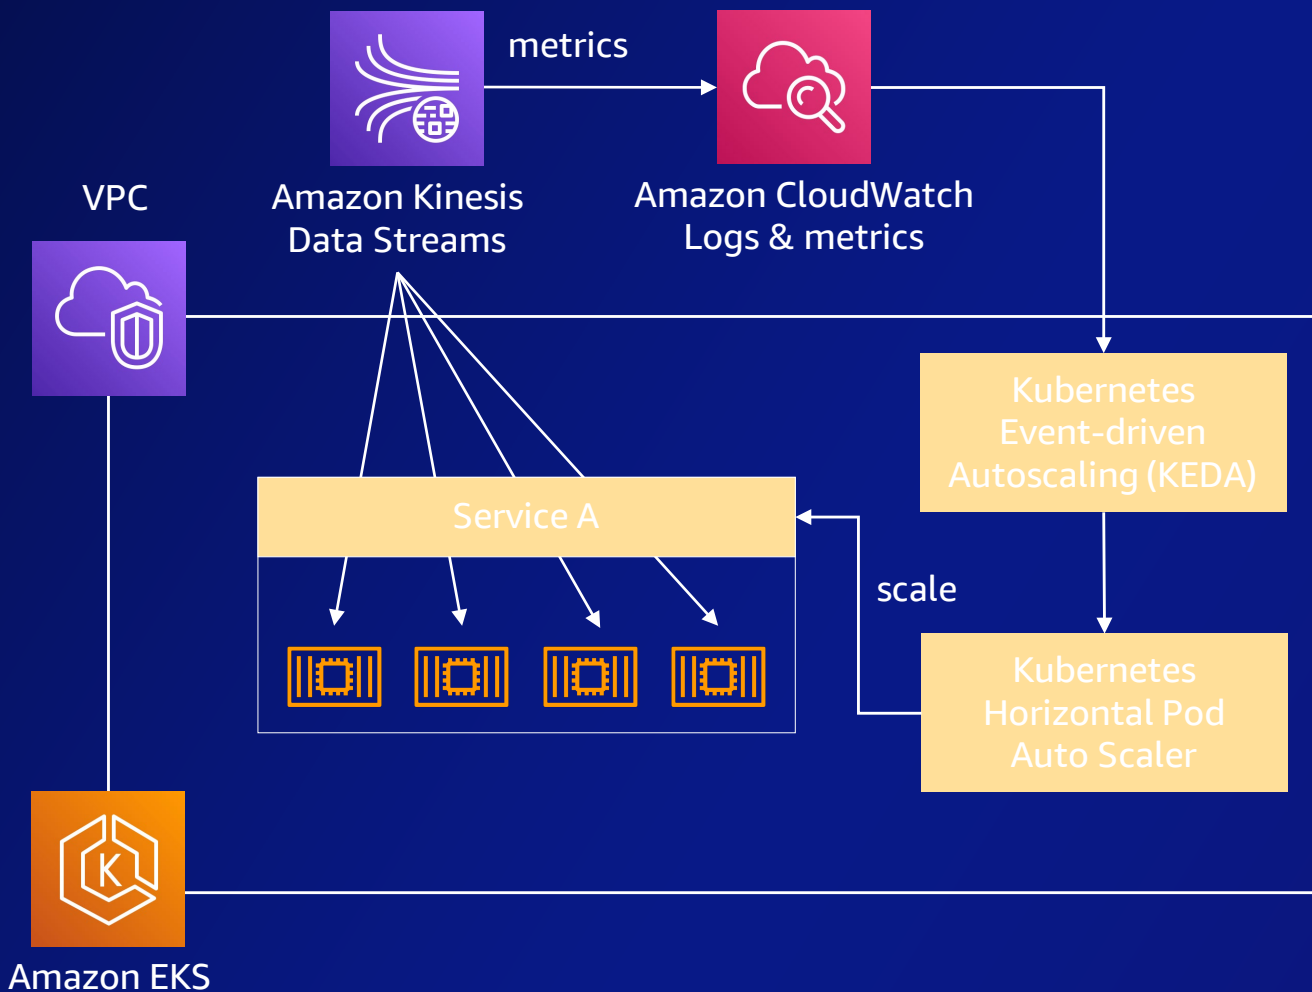
Event streaming pattern

Event streaming



- ❑ Kinesis Data Streams is used to collect and process large streams of data records in real time.
- ❑ Use API Gateway REST API as a Kinesis proxy to write data directly into a specified stream
- ❑ Kinesis Data Stream application (*consumer*) reads data from the streams using Kinesis Client Library (KCL)
- ❑ Consumer applications are run in ECS/EKS

Event streaming



- ❑ Kinesis Data Streams application (*consumer*) reads records from the stream using KCL
- ❑ Maximum consumers scale out is equal to the number of *shards* in the data stream
- ❑ Consumers may be scaled dynamically based on CloudWatch metrics for Kinesis Data Streams
- ❑ KEDA and HPA can be used to scale the number of pods

Data orchestration pattern

Use cases

File/video/image processing

Take collection of video files and convert them to sizes and formats that are ideal for different devices

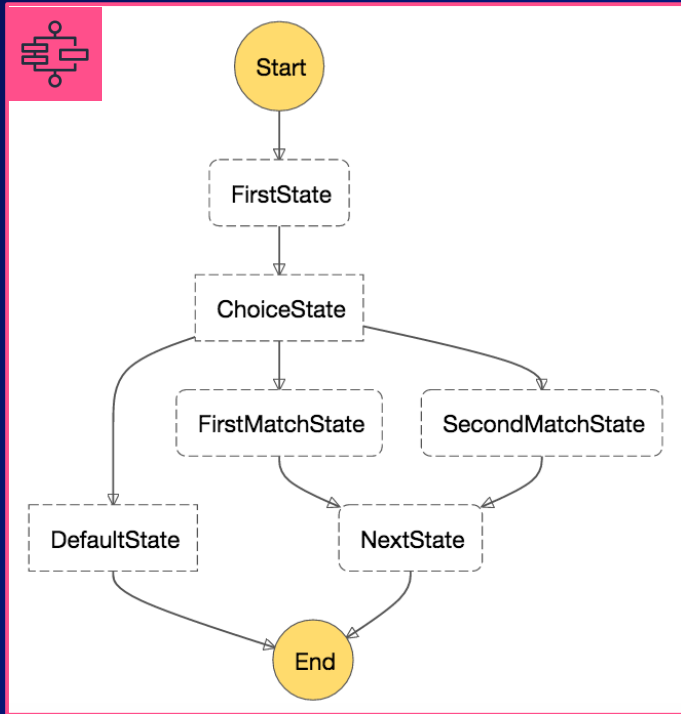
ETL/ELT pipelines

Combine sales opportunity records with marketing data to produce business intelligence report

Batch processing and HPC


Build a genomic secondary analysis pipeline that processes raw whole genome sequences into variant calls

Orchestrate with AWS Step Functions



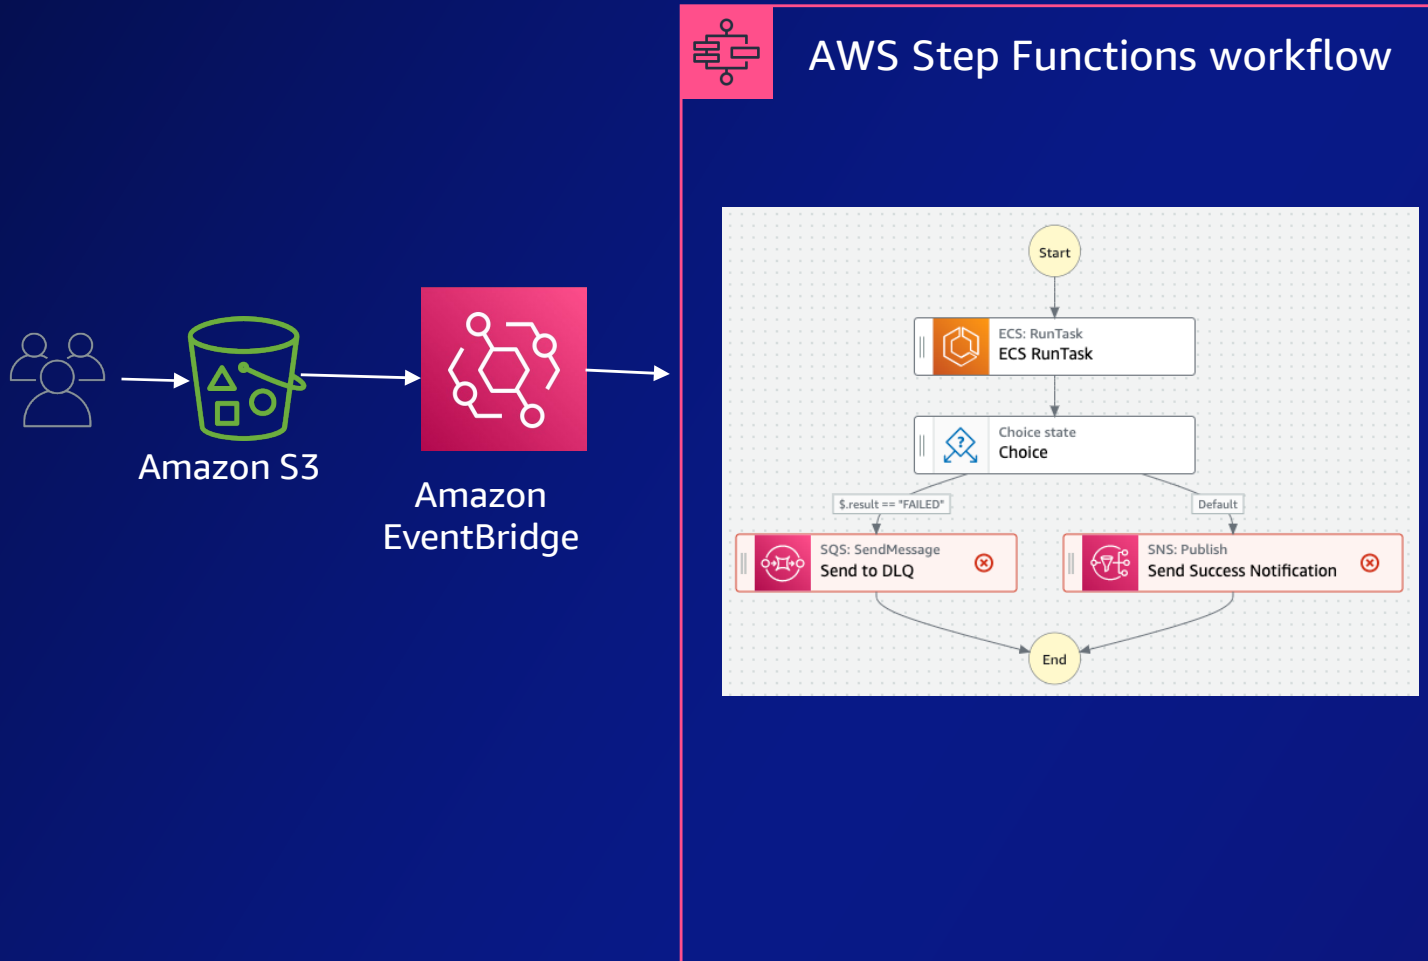
- ❑ Highly scalable and Serverless
- ❑ Low code orchestration
- ❑ Manages state of different activities
- ❑ Dependency management
- ❑ Integrates with 10k+ AWS API directly
- ❑ Built-in try/catch

Orchestrate on events from AWS



- ❑ Response to events from 110+ AWS services
- ❑ Example: Set Amazon S3 file notification for EventBridge when a new file is created

Orchestrate on events from AWS



- ❑ Response to events from 110+ AWS services
- ❑ Example: Set Amazon S3 file notification for EventBridge when a new file is created
- ❑ Set Rule in EventBridge EventBus to route to AWS Step Functions
- ❑ Build workflows with process the file, handle failures and success notifications
- ❑ 10k+ direct API calls including running task/pod in EKS and ECS

What next?



Start with existing patterns and modify

There are many sources for Cloud Native Architecture Patterns

- CDK Patterns - <https://cdkpatterns.com>
- Serverlessland Patterns Collection - <https://serverlessland.com/patterns>
- Construct Hub - <https://constructs.dev>
- AWS Compute Blog - <https://aws.amazon.com/blogs/compute/>
- AWS Samples - <https://github.com/aws-samples/serverless-patterns>
- Asynchronous Messaging Workshop - <https://github.com/aws-samples/asynchronous-messaging-workshop>

Create well-architected patterns

OPERATIONS

Understand the health and lifecycle of your application

RELIABILITY

Build resiliency and protect non-serverless resources

SECURITY

Focus on managing security boundaries and AppSec

PERFORMANCE

Optimize for low, steady, and/or peaks

COST

Factor in development, people, opportunity, and maintenance cost

Choose your pattern and go build!



skillbuilder.aws 

Your time is now

Build in-demand cloud skills *your way*

Thank you!

Mani Chandrasekaran

Principal Solutions Architect – India and South Asia

AWS India

manikach@amazon.com

<https://www.linkedin.com/in/cmanikandan>



Please complete the session survey

